

```
In [1]: #import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: data=pd.read_csv(r"C:\Users\user\Desktop\Vicky\4_drug200.csv")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age             200 non-null   int64
1   Sex             200 non-null   object
2   BP              200 non-null   object
3   Cholesterol     200 non-null   object
4   Na_to_K         200 non-null   float64
5   Drug            200 non-null   object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB
```

```
In [5]: data.shape
```

```
Out[5]: (200, 6)
```

```
In [6]: data.describe()
```

```
Out[6]:
```

	Age	Na_to_K
count	200.000000	200.000000
mean	44.315000	16.084485
std	16.544315	7.223956
min	15.000000	6.269000
25%	31.000000	10.445500
50%	45.000000	13.936500
75%	58.000000	19.380000
max	74.000000	38.247000


```
In [7]: data.columns
```

```
Out[7]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

```
In [9]: data1=data[['Age','Na_to_K']]
```


```
In [10]: sns.pairplot(data1)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x223b20ba9a0>
```



```
In [11]: sns.heatmap(data1.corr())
```

```
Out[11]: <AxesSubplot:>
```



```
In [12]: x=data[['Age', 'Na_to_K']]
         y=data1['Age']
```

```
In [13]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.1)
```

```
In [14]: from sklearn.linear_model import LinearRegression
         lr=LinearRegression()
         lr.fit(x_train,y_train)
```

```
Out[14]: LinearRegression()
```

```
In [15]: lr.intercept_
```

```
Out[15]: 7.105427357601002e-15
```


```
In [16]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=["Co-efficient"])
         coeff
```

```
Out[16]:
```

	Co-efficient
Age	1.000000e+00
Na_to_K	1.205361e-17

```
In [17]: prediction = lr.predict(x_train)
plt.scatter(y_train,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x223b44d8610>



```
In [18]: lr.score(x_test,y_test)
```

Out[18]: 1.0

```
In [19]: lr.score(x_train,y_train)
```

Out[19]: 1.0

```
In [20]: from sklearn.linear_model import Ridge,Lasso
```

```
In [21]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
```

Out[21]: 0.9999999586455789

```
In [22]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
la.score(x_test,y_test)
```

Out[22]: 0.9986414514540869

```
In [23]: from sklearn.linear_model import ElasticNet
en= ElasticNet()
en.fit(x_train,y_train)
```

Out[23]: ElasticNet()

```
In [24]: print(en.coef_)
```

[0.99632325 -0.]

```
In [26]: print(en.intercept_)
```

```
0.16271672981450536
```

```
In [27]: prediction = en.predict(x_test)  
prediction
```

```
Out[27]: array([20.08918168, 50.97520235, 26.06712116, 43.00461637, 34.03770714,  
                30.05241415, 59.94211157, 25.07079791, 73.89063704, 44.99726286,  
                67.91269755, 49.9788791 , 59.94211157, 34.03770714, 69.90534405,  
                46.98990936, 32.04506065, 31.0487374 , 68.9090208 , 28.05976766])
```

```
In [28]: print(en.score(x_test,y_train))
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-28-48d1f0543252> in <module>
----> 1 print(en.score(x_test,y_train))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py in score(self, X,
y, sample_weight)
    552         from .metrics import r2_score
    553         y_pred = self.predict(X)
--> 554         return r2_score(y, y_pred, sample_weight=sample_weight)
    555
    556     def _more_tags(self):

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in inn
er_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
---> 63             return f(*args, **kwargs)
    64
    65         # extra_args > 0

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in
r2_score(y_true, y_pred, sample_weight, multioutput)
    674         -3.0
    675         """
--> 676         y_type, y_true, y_pred, multioutput = _check_reg_targets(
    677             y_true, y_pred, multioutput)
    678         check_consistent_length(y_true, y_pred, sample_weight)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics\_regression.py in
_check_reg_targets(y_true, y_pred, multioutput, dtype)
    86         the dtype argument passed to check_array.
    87         """
---> 88         check_consistent_length(y_true, y_pred)
    89         y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    90         y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in che
ck_consistent_length(*arrays)
    260         uniques = np.unique(lengths)
    261         if len(uniques) > 1:
--> 262             raise ValueError("Found input variables with inconsistent num
bers of"
    263                               " samples: %r" % [int(l) for l in lengths])
    264

ValueError: Found input variables with inconsistent numbers of samples: [180,
20]
```

```
In [29]: from sklearn import metrics
```

```
In [30]: print("Mean Absolute error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute error: 0.05349675061312968

```
In [31]: print("Mean Absolute Square error:",metrics.mean_squared_error(y_test,predicti
```

Mean Absolute Square error: 0.003787627765347073

```
In [32]: print("Root mean Square error:",np.sqrt(metrics.mean_squared_error(y_test,pred
```

Root mean Square error: 0.061543706139190815

```
In [ ]:
```